Dpdk 环境搭建

工具准备： vmware + Ubuntu 16.04 Server 版本

1.vmware 添加两块网卡

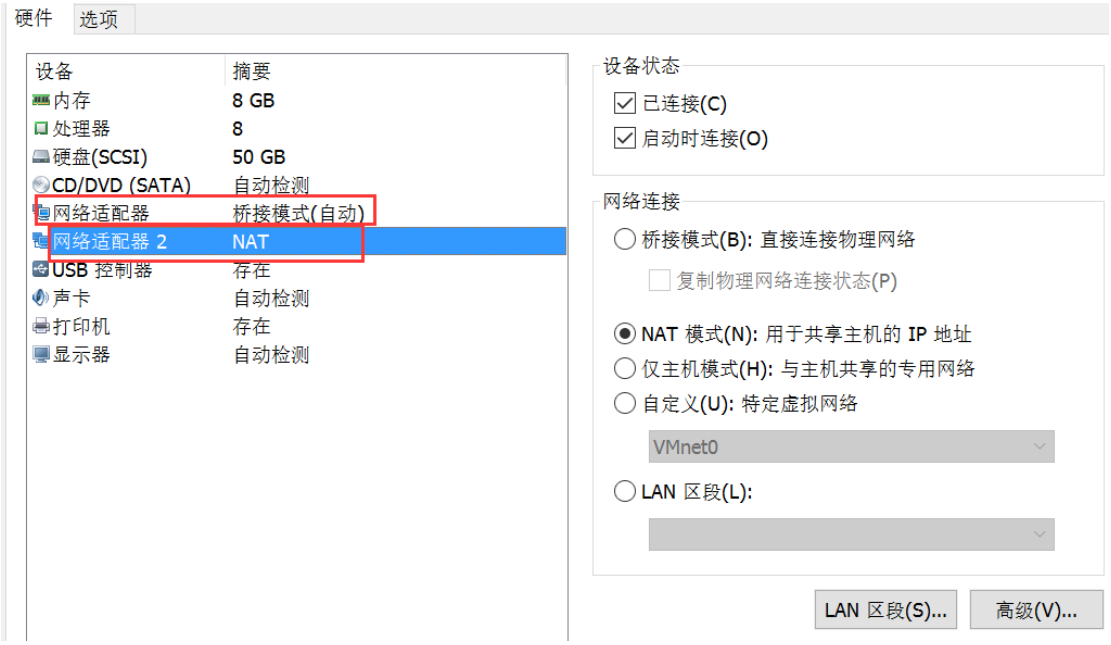

桥接网卡作为 DPDK 运行的网卡
NAT 网卡作为 ssh 连接的网卡

2.修改网卡配置信息

```
sata0:1.deviceType = "cdrom-raw"
ethernet0.present = "TRUE"
ethernet0.connectionType = "bridged"
ethernet0.virtualDev = "vmxnet3"
ethernet0.wakeOnPcktRcv = "TRUE"
ethernet0.addressType = "generated"
usb.present = "TRUE"
ehci.present = "TRUE"
ehci.pciSlotNumber = "35"
sound.present = "TRUE"
sound.fileName = "-1"
```

由e1000修改为

将 ethernet0.virtualDev 由 e1000 修改 vmxnet3，因为 vmware 的 vmxnet3 支持多队列网卡

3.修改 ubuntu 系统的启动参数

物理机：

```
default_hugepages=1G hugepagesz=1G hugepages=20 isolcpus=0-7
```

虚拟机：

```
default_hugepages=1G hugepagesz=2M hugepages=1024 isolcpus=0-2
```

4.查看系统是否支持多队列网卡

执行：cat /proc/interrupts

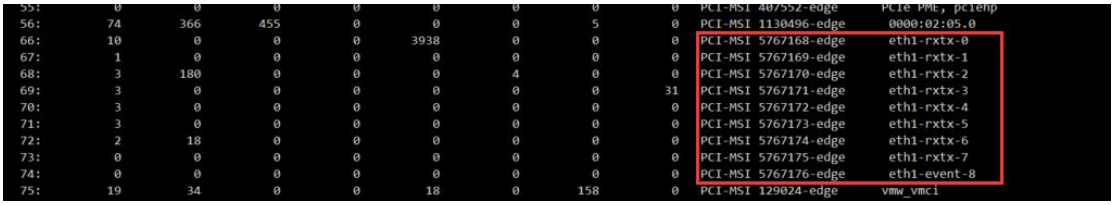

5.编译 DPDK

下载 dpdk ： https://core.dpdk.org/download/

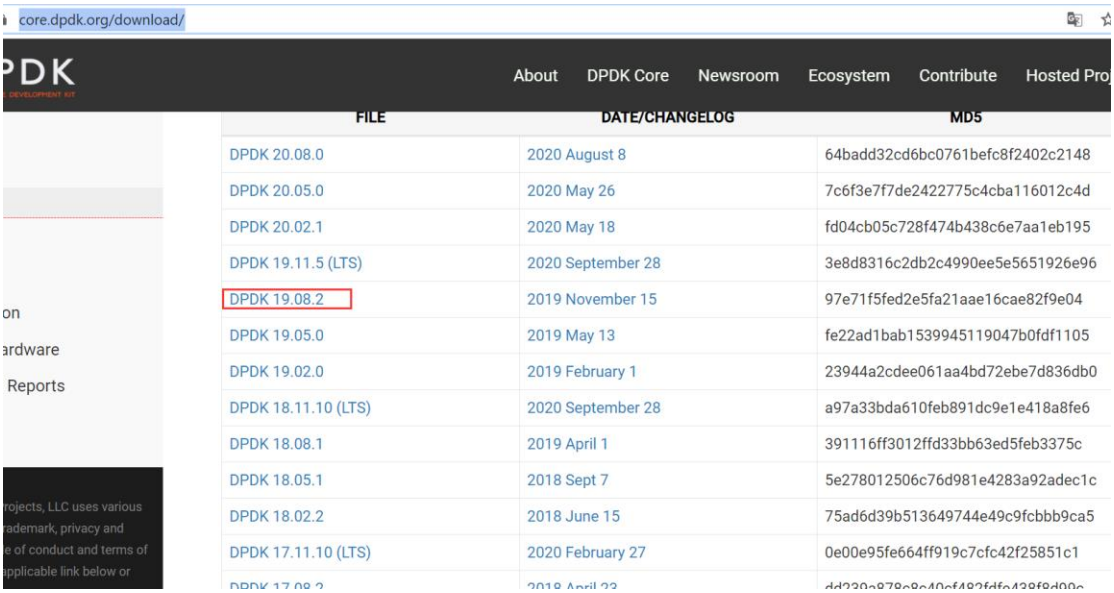

随便挑选版本，不同版本直接子系统接口会有差异，建议选择 dpdk 19.08.2

可以通过 usertools/dpdk-setup.sh

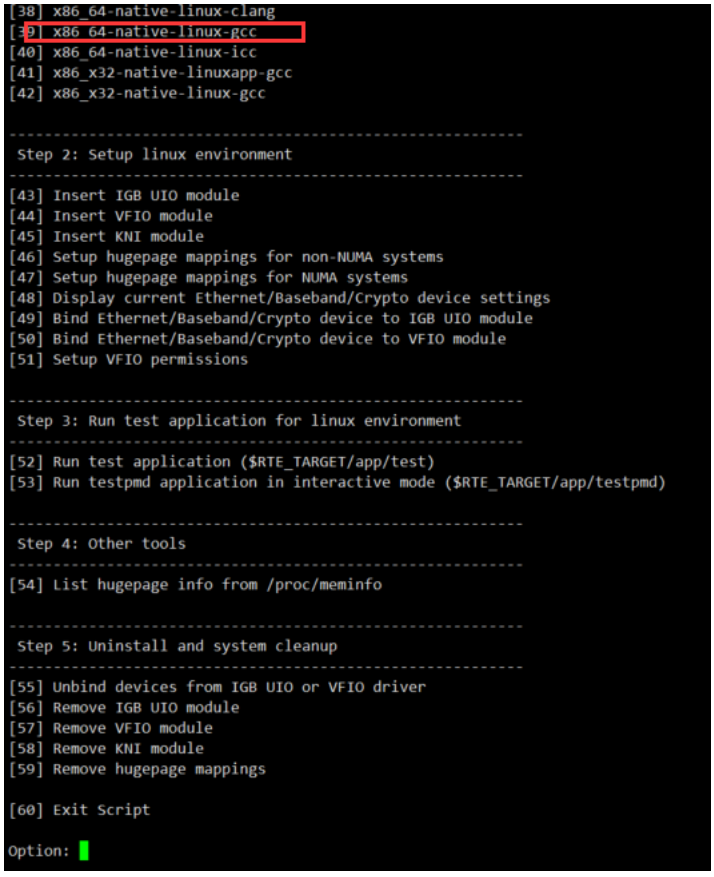

64 位系统选择 39.

编译完会多出 x86_64-native-linux-gcc 的文件夹

6.设置 DPDK 的环境变量

```
# export RTE_SDK=/home/dpdk
# export RTE_TARGET=x86_64-native-linux-gcc
```

7.执行 testpmd

执行

```
# /usertools/dpdk-setup.sh
```

```
------------------------------------------------------------
 Step 2: Setup linux environment
------------------------------------------------------------
[43] Insert IGB UIO module
[44] Insert VFIO module
[45] Insert KNI module
[46] Setup hugepage mappings for non-NUMA systems
[47] Setup hugepage mappings for NUMA systems
[48] Display current Ethernet/Baseband/Crypto device settings
[49] Bind Ethernet/Baseband/Crypto device to IGB UIO module
[50] Bind Ethernet/Baseband/Crypto device to VFIO module
[51] Setup VFIO permissions

------------------------------------------------------------
 Step 3: Run test application for linux environment
------------------------------------------------------------
[52] Run test application ($RTE_TARGET/app/test)
[53] Run testpmd application in interactive mode ($RTE_TARGET/app/testpmd)

------------------------------------------------------------
 Step 4: Other tools
------------------------------------------------------------
[54] List hugepage info from /proc/meminfo

------------------------------------------------------------
 Step 5: Uninstall and system cleanup
------------------------------------------------------------
[55] Unbind devices from IGB UIO or VFIO driver
[56] Remove IGB UIO module
[57] Remove VFIO module
[58] Remove KNI module
[59] Remove hugepage mappings

[60] Exit Script

Option:
```

选择 43 插入 IGB_UIO 模块， 选择网卡为 vmxnet3 会加载此模块
选择 44 插入 VFIO 模块，选择网卡为 e1000 会加载此模块

选择 49 绑定 igb_uio 模块， 也可以退出，通过命令来执行。

```
# ifconfig eth0 down
# /usertools/dpdk-devbind.py --bind=igb_uio eth0
```

选择 53 运行 testpmd

```
Option: 53


  Enter hex bitmask of cores to execute testpmd app on
  Example: to execute app on cores 0 to 7, enter 0xff
bitmask: 7
Launching app
EAL: Detected 8 lcore(s)
EAL: Detected 1 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'PA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
EAL: PCI device 0000:03:00.0 on NUMA socket -1
EAL:   Invalid NUMA socket, default to 0
EAL:   probe driver: 15ad:7b0 net_vmxnet3
EAL: PCI device 0000:0b:00.0 on NUMA socket -1
EAL:   Invalid NUMA socket, default to 0
EAL:   probe driver: 15ad:7b0 net_vmxnet3
Interactive-mode selected
testpmd: create a new mbuf pool <mbuf_pool_socket_0>: n=163456, size=2176, socket=0
testpmd: preferred mempool ops selected: ring_mp_mc

Warning! port-topology=paired and odd forward ports number, the last port will pair with itself.

Configuring Port 0 (socket 0)
Port 0: 00:0C:29:18:EF:9D
Checking link statuses...
Done
testpmd>
```

```
> show port info 0
```

```
testpmd> show port info 0

********************* Infos for port 0  *********************
MAC address: 00:0C:29:18:EF:9D
Device name: 0000:03:00.0
Driver name: net_vmxnet3
Connect to socket: 0
memory allocation on the socket: 0
Link status: up
Link speed: 10000 Mbps
Link duplex: full-duplex
MTU: 1500
Promiscuous mode: enabled
Allmulticast mode: disabled
Maximum number of MAC addresses: 1
Maximum number of MAC addresses of hash filtering: 0
VLAN offload:
  strip off
  filter off
  qinq(extend) off
Supported RSS offload flow types:
  ipv4
  ipv4-tcp
  ipv6
  ipv6-tcp
Minimum size of RX buffer: 1646
Maximum configurable length of RX packet: 16384
Current number of RX queues: 1
Max possible RX queues: 16
Max possible number of RXDs per queue: 4096
Min possible number of RXDs per queue: 128
RXDs number alignment: 1
Current number of TX queues: 1
Max possible TX queues: 8
Max possible number of TXDs per queue: 4096
Min possible number of TXDs per queue: 512
TXDs number alignment: 1
Max segment number per packet: 255
Max segment number per MTU/TSO: 16
testpmd>
```

8.编译 DPDK 程序

进入 example/helloworld ，
可以直接 make，
也可以通过 gcc 命令编译

```
# gcc -o helloword main.c -I /usr/local/include/dpdk/ -ldpdk -lpthread -lnuma -ldl
```

```
root@ubuntu:/home/king/share/dpdk/dpdk-stable-19.08.2/examples/helloworld# gcc -o helloword main.c -I /usr/local/include/dpdk/ -ldpdk -lpthread -lnuma -ldl
root@ubuntu:/home/king/share/dpdk/dpdk-stable-19.08.2/examples/helloworld#
root@ubuntu:/home/king/share/dpdk/dpdk-stable-19.08.2/examples/helloworld#
root@ubuntu:/home/king/share/dpdk/dpdk-stable-19.08.2/examples/helloworld#
root@ubuntu:/home/king/share/dpdk/dpdk-stable-19.08.2/examples/helloworld# ./helloword -l 0-3 -n 4
EAL: Detected 8 lcore(s)
EAL: Detected 1 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'PA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
hello from core 1
hello from core 2
hello from core 3
hello from core 0
root@ubuntu:/home/king/share/dpdk/dpdk-stable-19.08.2/examples/helloworld#
```

9. 运行 DPDK 案例

Kni 运行：

```
./build/kni -l 4-7 -n 4 -- -P -p 0x3 -m --config="(0, 4, 6),(1, 5, 7)"
```

L3fwd 运行

```
# ./build/l3fwd -l 4-7 -n 4 -- -p 0x3 --config="(0, 0, 4),(1, 0, 5)" --parse-ptype
```

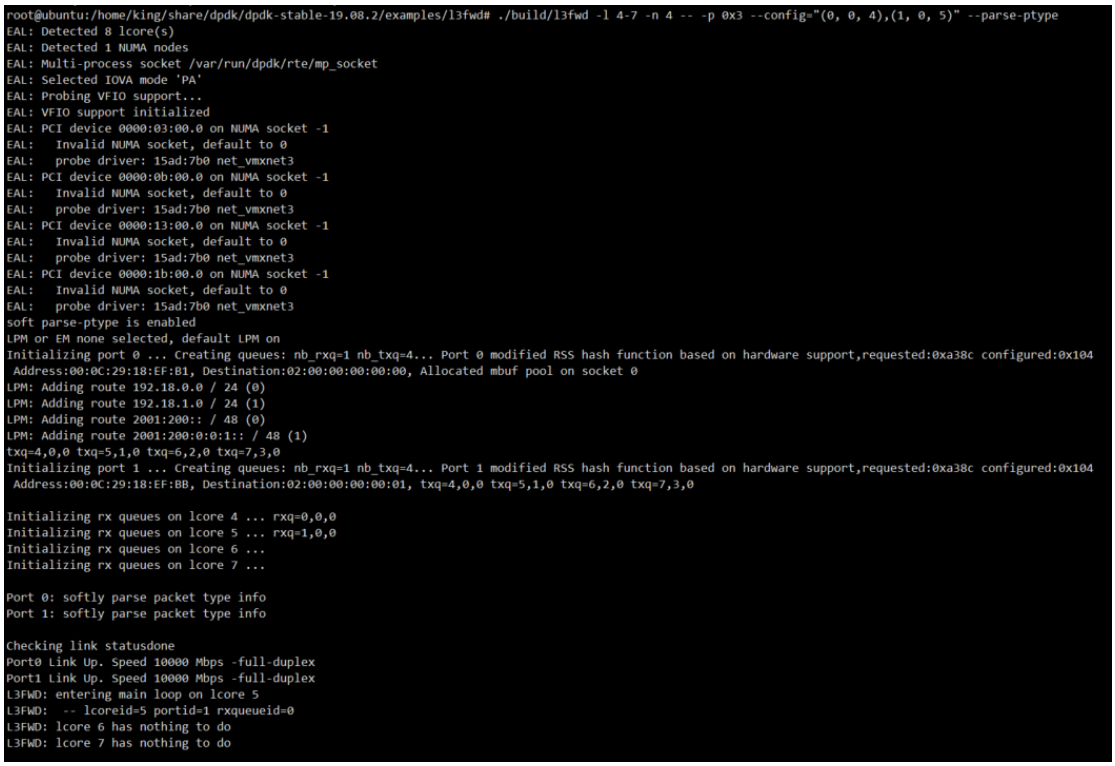

常见错误

错误 1：

没有设置环境变量，请查看本篇《设置 dpdk 的环境变量》

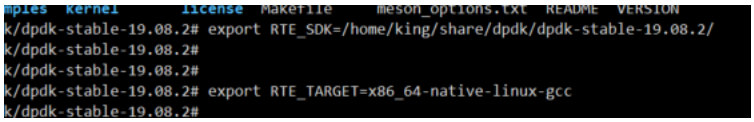

```
  Enter hex bitmask of cores to execute testpmd app on
  Example: to execute app on cores 0 to 7, enter 0xff
bitmask: 4
Launching app
EAL: Detected 8 lcore(s)
EAL: Detected 1 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'PA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
EAL: PCI device 0000:03:00.0 on NUMA socket -1
EAL:   Invalid NUMA socket, default to 0
EAL:   probe driver: 15ad:7b0 net_vmxnet3
EAL: PCI device 0000:0b:00.0 on NUMA socket -1
EAL:   Invalid NUMA socket, default to 0
EAL:   probe driver: 15ad:7b0 net_vmxnet3
EAL: Error - exiting with code: 1
  Cause: No cores defined for forwarding
Check the core mask argument
```

bitmask 请选择 7, bit:111